

4. Threads

ECE30021/ITP30002 Operating Systems

Agenda



- Overview
- Multithreading models
- Thread libraries
- Threading issues
- Operating system examples

Motivation



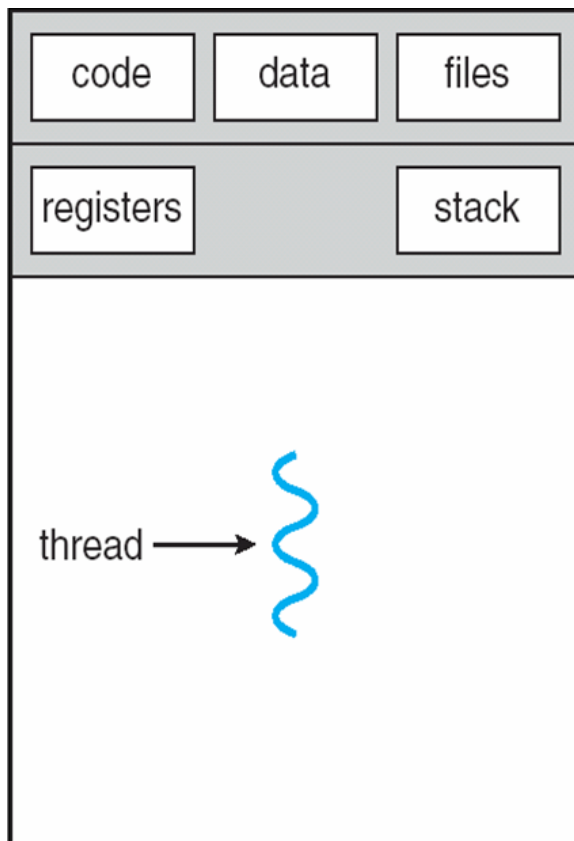
- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
 - Update display, fetch data, spell checking, answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

Overview

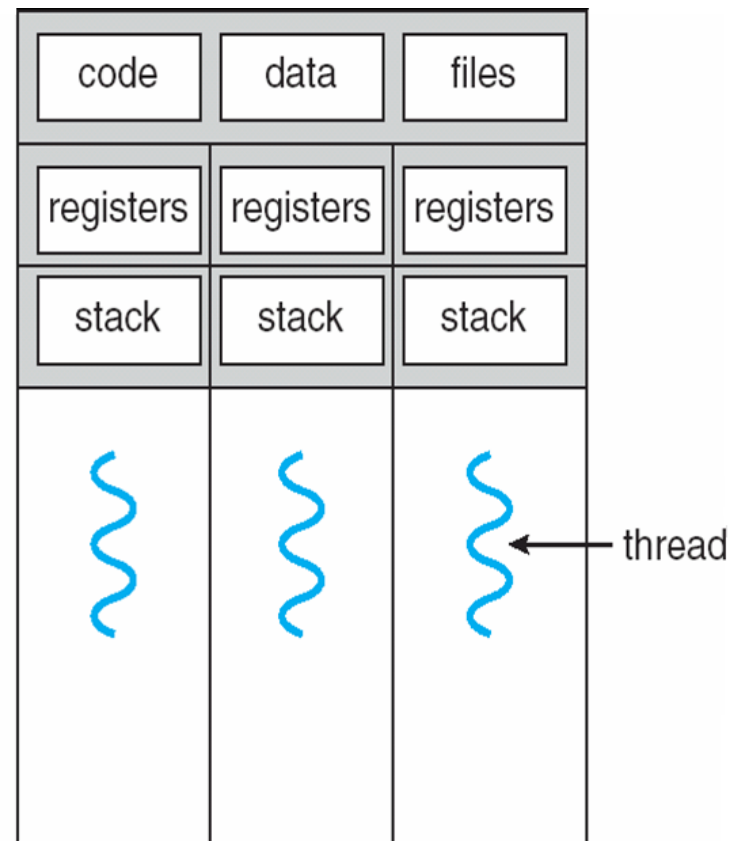


- **Process**: program in execution
 - Each process occupies resources required for execution
- **Thread**: a way for a program to split itself into two or more simultaneously running tasks
 - Smaller unit than process
 - Threads in a process share resources
- A thread is comprised of
 - Thread ID, program counter, register set, stack, etc.

Multithreaded Process

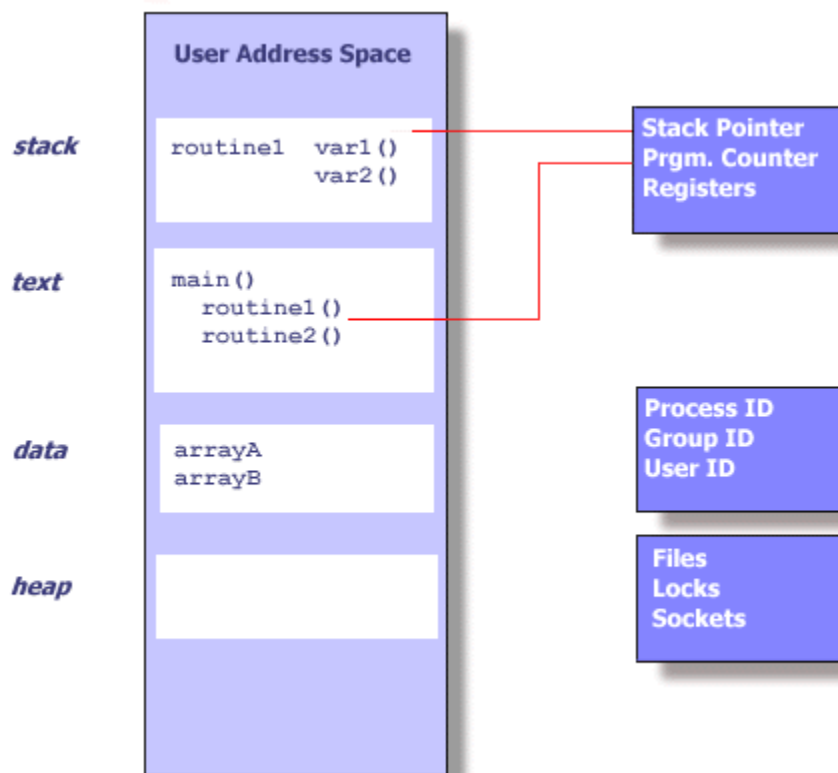


single-threaded process

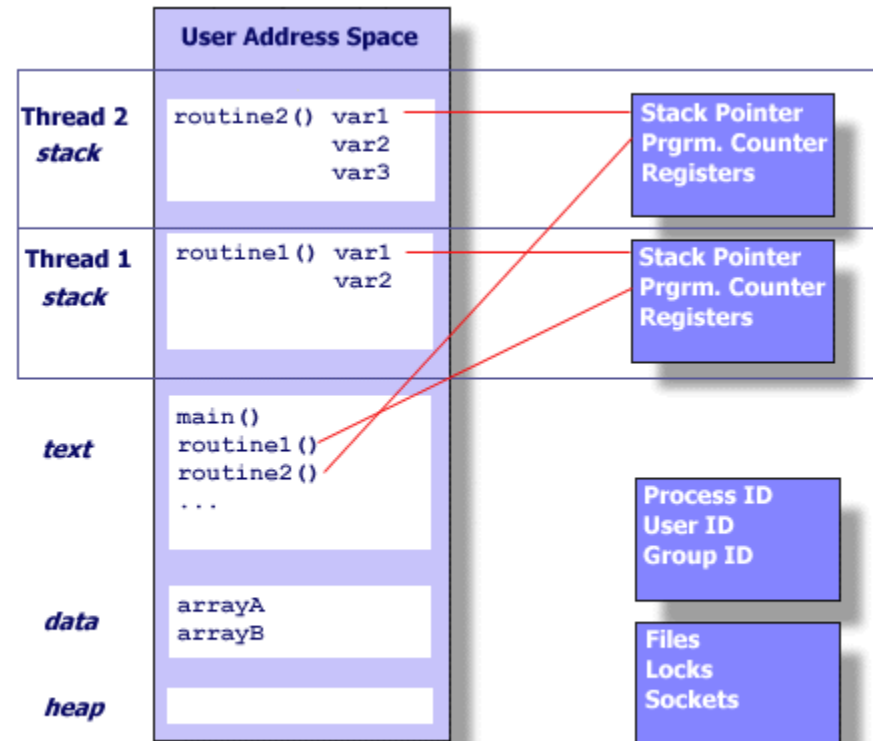


multithreaded process

Process vs. Thread



A process in UNIX

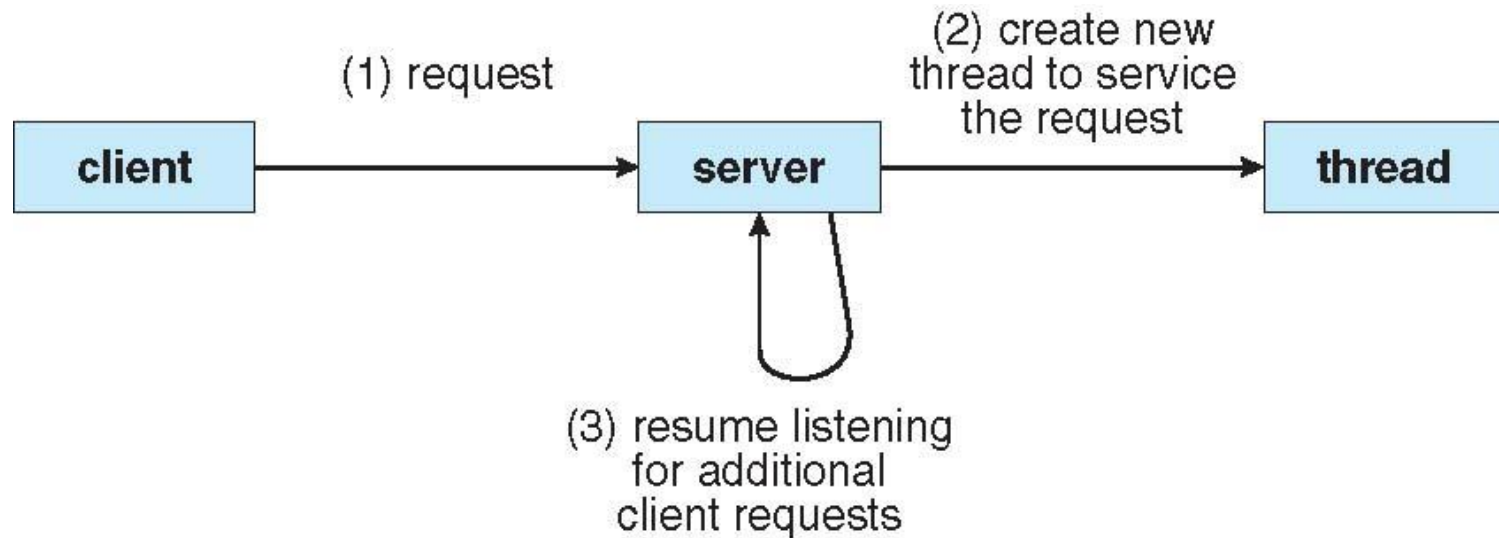


Threads within a Process

Why Thread?

- Process creation is expensive in time and resource

Ex) Web server accepting thousands of requests



Why Thread?



- Advantages of multithreaded programming
 - Responsiveness
 - Resource sharing
 - Economy
 - Creating process is about 30 times slower than creating thread
 - Scalability
 - Utilization of multiprocessor architectures

Multicore Programming



- Multicore or multiprocessor systems putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging

Multicore Programming (Cont.)



■ Types of parallelism

- Data parallelism – distributes subsets of the same data across multiple cores, same operation on each
- Task parallelism – distributing threads across cores, each thread performing unique operation

■ As # of threads grows, so does architectural support for threading


- CPUs have cores as well as hardware threads
- Consider Oracle SPARC T4 with 8 cores, and 8 hardware threads per core

Concurrency vs. Parallelism

- **Concurrency** supports more than one task making progress


- Single processor / core, scheduler providing concurrency

Ex) Concurrent execution on single-core system:



- **Parallelism** implies a system can perform more than one task simultaneously

Ex) Parallelism on a multi-core system:



Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components
 - S is serial portion, N processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

Ex) Application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times


- As N approaches infinity, speedup approaches 1 / S
- **Serial portion of an application has disproportionate effect on performance gained by adding additional cores.**

Types of Thread



- **User thread:** thread supported by **thread library** in user level
 - Created by library function call (not system call)
 - Kernel is not concerned in user thread.
 - Switching of user thread is faster than kernel thread.
- **Kernel thread:** thread supported by **kernel**
 - Created and managed by kernel
 - Scheduled by kernel
 - Cheaper than process
 - More expensive than user thread

Kernel Thread




Agenda



- Overview
- Multithreading models
- Thread libraries
- Threading issues
- Operating system examples

Multithreading Models

- Major issue: correspondence between user threads and kernel threads




Multithreading Models

■ Many-to-one model

- Many user threads are mapped to single kernel thread
- Threads are managed by user-level thread library


Ex) Green threads, GNU Portable Threads



■ One-to-one model

- Each user thread is mapped to a kernel thread
- Provides more concurrency
- Problem: overhead


Ex) Linux, Windows, Solaris



Multithreading Models


■ Many-to-Many model

- Multiplex many user level threads to smaller or equal number of kernel threads
- Compromise between $n:1$ model and $1:1$ model



■ Two-level model

- Variation of $N:M$ model
 - Basically $N:M$ model
 - A user thread can be bound to a kernel thread
- Ex) IRIX, HP-UX, Tru64, Solaris (old)




Scheduler Activation



- An issue about many-to-many model and two-level model: communication between the kernel and the thread library.
- **Scheduler activation**: one scheme for communication between user thread library and kernel
- In many-to-many model and two-level model, user threads are connected with kernel threads through **LWP**
- **Lightweight process (LWP)**
 - A data structure connecting user thread to kernel thread.
 - Basically, a LWP corresponds to a kernel thread, but there are some exceptions.
 - To the user-level thread library, a LWP appears to be a **virtual processor**.

Scheduler Activation and LWP

- Connection between user / kernel threads through LWP



Scheduler Activation



■ Scheduler activation

- Kernel provides a set of virtual processors(LWP's).
- User level thread library schedules user threads onto virtual processors.
- If a kernel thread is blocked or unblocked, kernel notices it to thread library(**upcall**).
- Upcall handler schedules properly.
 - If a kernel thread is blocked, assign the LWP to another thread
 - If a kernel thread is unblocked, assign an LWP to it.

Agenda



- Overview
- Multithreading models
- Thread libraries
- Threading issues
- Operating system examples

Thread Libraries



- **Thread library:** set of API's to create and manage threads
 - User level library
 - Kernel level library

Examples)

- POSIX Pthreads: a specification with various implementations.
 - LinuxThreads
 - NPTL (Native POSIX Thread Library)
 - GNU Portable Threads
 - Open source Pthreads for win32
 - Etc.
- Win32 threads
- Java threads

POSIX Pthreads



- Reading assignment: Search the Internet for the following functions, and study them.
- API functions
 - `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void * (*start_routine)(void *), void *arg);`
 - Creates thread
 - `int pthread_attr_init(pthread_attr_t *attr);`
 - Initializes *attr* by default values
 - `int pthread_join(pthread_t th, void **thread_return);`
 - Waits for a thread *th*.
 - `void pthread_exit(void *retval);`
 - Terminates thread

Example



```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* the thread */

int main (int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */
    if (argc < 2 ) {
        fprintf(stderr, "usage: a.out <integer>\n");
        exit(0);
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be <= 0\n", atoi(argv[1]));
        exit(0);
    }

    /* get the default attributes */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* now wait for the thread to exit */
    pthread_join(tid, NULL);
    printf("sum = %d\n", sum);
    return 0;
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int upper = atoi (param);
    int i;
    sum = 0;
    if (upper > 0) {
        for (i = 1; i <= upper ; i++)
            sum += i;
    }
    pthread_exit(0);
}
```

Win32 Threads

- Create

```
HANDLE WINAPI CreateThread(  
    LPSECURITY_ATTRIBUTES lpThreadAttributes,  
    SIZE_T dwStackSize,  
    LPTHREAD_START_ROUTINE lpStartAddress,  
    LPVOID lpParameter,  
    DWORD dwCreationFlags,  
    LPDWORD lpThreadId  
);
```

See <http://msdn.microsoft.com/library/default.asp?url=/library/en-us/dllproc/base/createthread.asp>

- Wait

```
DWORD WINAPI WaitForSingleObject(  
    HANDLE hHandle,  
    DWORD dwMilliseconds  
);
```

- Close (deallocate) handle

```
BOOL CloseHandle(LPDWORD lpThreadId);
```

Windows Multithreaded C Program

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }

    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    if (ThreadHandle != NULL) {
        /* now wait for the thread to finish */
        WaitForSingleObject(ThreadHandle, INFINITE);

        /* close the thread handle */
        CloseHandle(ThreadHandle);

        printf("sum = %d\n", Sum);
    }
}
```

Java Threads



- Java threads are managed by the JVM
 - Typically, implemented using the threads model provided by underlying OS
- Java threads may be created by:
 - Extending **Thread** class
 - For detail, search 'Java Thread class' from Internet
 - Implementing the **Runnable** interface

```
public interface Runnable
{
    public abstract void run();
}
```

Java Thread using Thread Class

■ Extending Thread class

```
class PrimeThread extends Thread {  
    long minPrime;  
    PrimeThread(long minPrime) {  
        this.minPrime = minPrime;  
    }  
  
    public void run() {  
        // compute primes larger than minPrime  
        . . .  
    }  
}
```

■ Launching thread

```
PrimeThread p = new PrimeThread(143);  
p.start();
```

Java Thread using Running Interface



■ Extending Thread class

```
class PrimeRun implements Runnable {  
    long minPrime;  
    PrimeRun(long minPrime) {  
        this.minPrime = minPrime;  
    }  
  
    public void run() {  
        // compute primes larger than minPrime  
        ...  
    }  
}
```

■ Launching thread

```
PrimeRun p = new PrimeRun(143);  
new Thread(p).start();
```

Java Multithreaded Program

```
class Sum
{
    private int sum;

    public int getSum() {
        return sum;
    }

    public void setSum(int sum) {
        this.sum = sum;
    }
}

class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }

    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setSum(sum);
    }
}
```

Java Multithreaded Program (Cont.)

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                Sum sumObject = new Sum();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sumObject));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sumObject.getSum());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>"); }
}
```


Implicit Threading



- Creation and management of threads done by **compilers** and **run-time libraries** rather than programmers
- Three methods explored
 - Thread Pools
 - OpenMP
 - Grand Central Dispatch
- Other methods include Microsoft Threading Building Blocks (TBB), `java.util.concurrent` package

Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
 - Separating task to be performed from mechanics of creating task allows different strategies for running task
 - i.e.Tasks could be scheduled to run periodically
- Windows API supports thread pools:

```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
    * this function runs as a separate thread.  
    */  
}
```

OpenMP



- Provides support for parallel programming in shared-memory environments
 - Set of compiler directives and an API for C, C++, FORTRAN
 - Identifies parallel regions – blocks of code that can run in parallel
- Create as many threads as there are cores

```
#pragma omp parallel
printf("Hello, World!\n");
```
- Run for loop in parallel

```
#pragma omp parallel for
for(i=0;i<N;i++) {
    c[i] = a[i] + b[i];
}
```

Grand Central Dispatch



- Apple technology for Mac OS X and iOS operating systems
- Extensions to C, C++ languages, API, and run-time library
- Allows identification of parallel sections
- Manages most of the details of threading
- Block is in “`^{ }`” – `^{ printf("I am a block"); }`
- Blocks placed in dispatch queue
 - Assigned to available thread in thread pool when removed from queue

Grand Central Dispatch



■ Two types of dispatch queues:

- serial – blocks removed in FIFO order, queue is per process, called main queue
 - Programmers can create additional serial queues within program
- concurrent – removed in FIFO order but several may be removed at a time
 - Three system wide queues with priorities low, default, high

```
dispatch_queue_t queue = dispatch_get_global_queue  
    (DISPATCH_QUEUE_PRIORITY_DEFAULT, 0);  
  
dispatch_async(queue, ^{ printf("I am a block."); });
```

Agenda



- Overview
- Multithreading models
- Thread libraries
- Threading issues
- Operating system examples

Threading Issues



- `fork()` and `exec()`
- Cancellation
- Signal handling
- Thread-local storage
- Scheduler activation (already covered)

fork() and exec()



- fork() on multithreaded process
 - Duplicates all threads in the process?
 - Duplicates only corresponding thread?
 - UNIX supports two versions of fork
 - fork(), fork1()
- exec() on multithreaded process
 - Replace entire process

Cancellation

■ Thread cancellation

- Terminating a thread (target thread) before it has completed

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);
```

■ Problem with thread cancellation

- A thread share the resource with other threads
cf. A process has its own resource.
→ A thread can be cancelled while it updates data shared with other threads

Cancellation



- Cancellation scenarios
 - Asynchronous cancellation
 - Immediate termination
 - Deferred cancellation
 - Target thread checks periodically whether it should terminate.
Ex) `pthread_testcancel()` of `pthread`
 - Safer than asynchronous cancellation

Thread Cancellation (Cont.)

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is deferred
 - Cancellation only occurs when thread reaches cancellation point
 - I.e. `pthread_testcancel()`
 - Then cleanup handler is invoked
- On Linux systems, thread cancellation is handled through signals

Signal Handling



- **Signal:** mechanism provided by UNIX to notify a process a particular event has occurred
 - A signal can be generated by various sources.
 - The signal is delivered to **a process**.
 - The process handles it.
 - Default signal handler (kernel)
 - User-defined signal handler
- **Types of signal**
 - Synchronous: signal from same process
Ex) illegal memory access, division by 0
 - Asynchronous: signal from external sources
Ex) <Ctrl>-C

What is a Signal?



- A signal is a software interrupt delivered to a process.
 - The operating system uses signals to report exceptional situations to an executing program
- A signal is a limited form of inter-process communication used in Unix and other POSIX-compliant operating systems.
 - Essentially it is an asynchronous notification sent to a process in order to notify it of an event that occurred

Examples of UNIX Signals

■ Signals in UNIX (in signal.h)

```
#define SIGHUP      1  /* hangup */
#define SIGINT      2  /* interrupt */
#define SIGQUIT     3  /* quit */
#define SIGILL      4  /* illegal instruction (not reset when caught) */
#define SIGTRAP     5  /* trace trap (not reset when caught) */
#define SIGIOT      6  /* IOT instruction */
#define SIGABRT     6  /* used by abort, replace SIGIOT in the future */
#define SIGEMT      7  /* EMT instruction */
#define SIGFPE      8  /* floating point exception */
#define SIGKILL     9  /* kill (cannot be caught or ignored) */
#define SIGBUS     10  /* bus error */
#define SIGSEGV     11  /* segmentation violation */
#define SIGSYS      12  /* bad argument to system call */
#define SIGPIPE     13  /* write on a pipe with no one to read it */
#define SIGALRM     14  /* alarm clock */
#define SIGTERM     15  /* software termination signal from kill */
```

Installing Signal Handler

■ Defining signal handler

```
typedef void (*sighandler_t)(int);  
sighandler_t signal(int signum,  
                    sighandler_t handler);
```

■ Example

```
#include <signal.h>  
#include <unistd.h>  
  
void sig_handler(int signo);  
  
// press CTRL-W to terminate this program  
int main()  
{  
    int i = 0;  
    signal(SIGINT, (void *)sig_handler);  
  
    while(1){  
        printf("%dWn", i++);  
        sleep(1);  
    }  
    return 1;  
}  
  
void sig_handler(int signo)  
{  
    printf("SIGINT was received!Wn");  
}
```

Signal Handling



- Question: To what thread the signal should be delivered?
- Possible options
 - To the thread to which the signal applies
 - To every thread in the process
 - To certain threads in the process
 - Assign a specific thread to receive all signals
 - ➔ depend on type of signal
- Another scheme: specify a thread to deliver the signal
Ex) `pthread_kill(tid, signal)` in POSIX

Thread-Local Storage



- **Thread-local storage** (TLS) allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- Different from local variables
 - Local variables visible only during single function invocation
 - TLS visible across function invocations
- Similar to static data
 - TLS is unique to each thread

Agenda



- Overview
- Multithreading models
- Thread libraries
- Threading issues
- Operating system examples

Windows XP Threads



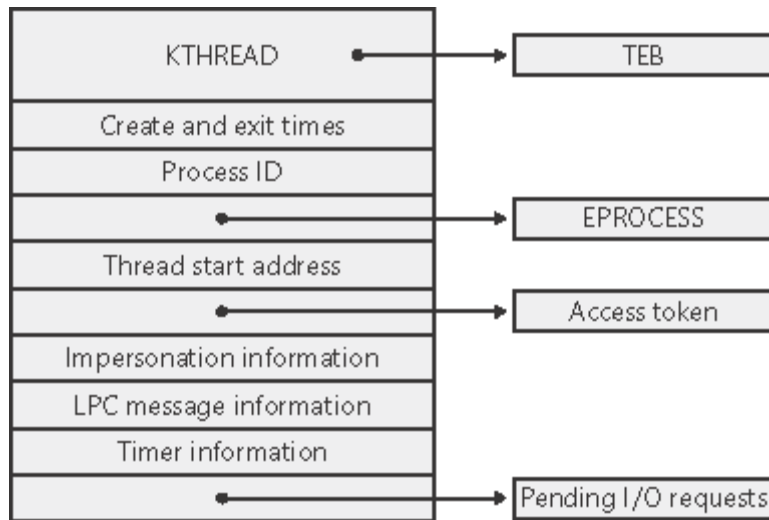
- Implements the one-to-one mapping, kernel-level
 - Additionally, many-to-many model is supported by fiber library
 - Each thread contains
 - A thread id, register set, separate user and kernel stacks, private data storage area (for run-time libraries and DLLs)
- Cf. The register set, stacks, and private storage area are known as **the context of the thread**.

Internal Data Structures

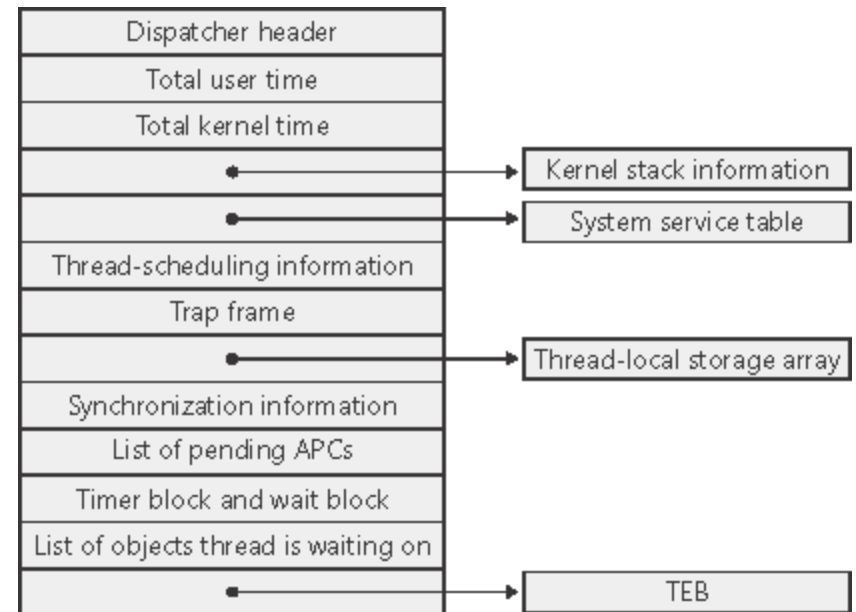


- The primary data structures of a thread include:
 - ETHREAD (executive thread block)
 - KTHREAD (kernel thread block)
 - Information for thread scheduling and synchronization
 - TEB (thread environment block)
 - Context information for the image loader and various Windows DLL

Layout of ETHREAD and KTHREAD

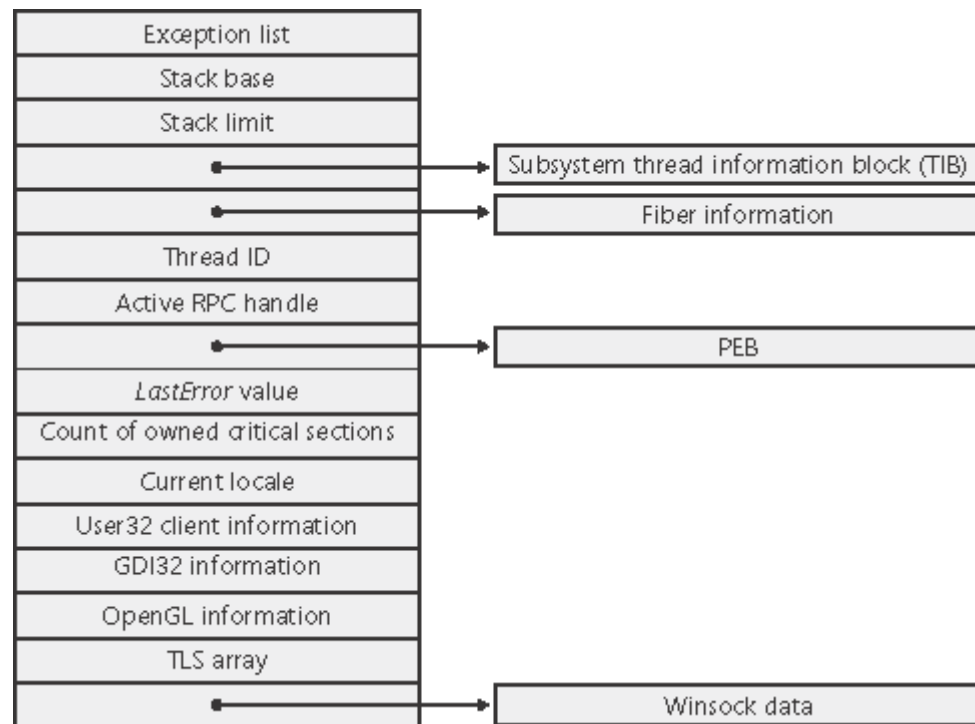


ETHREAD

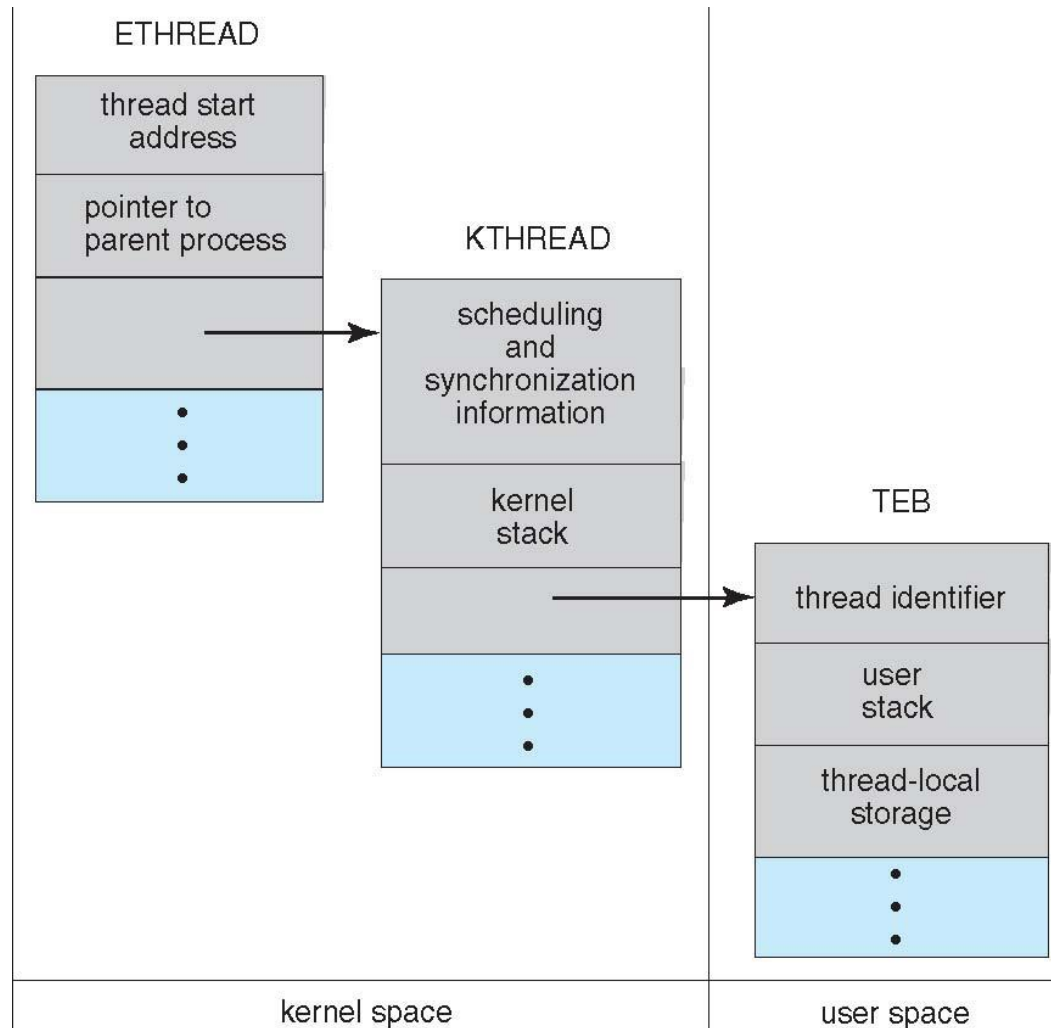


KTHREAD

Layout of TEB



Windows XP Threads



Linux Threads



- First implementation: LinuxThreads
 - Thread is created by `clone()` system call
 - However, no difference between process and thread
 - Term '**task**' is used more frequently than 'process' or 'thread'
 - User can control how much the resource is shared between parent and child.

Linux Threads



■ Problems of *LinuxThreads*

- It did not scale well,
 - There is a limitation on the number of threads, generally between 1024 and 8192.
 - The overhead of creating and destroying processes is relatively high.
- The manager thread resulted in some fragility and another scaling bottleneck.
- Some required POSIX semantics were not possible. Each thread had its own PID.

■ NPTL (Native POSIX Thread Library)

- Better than LinuxThreads in performance and compatibility